Fundamental Concepts in Computational and Applied Mathematics

Juan Meza
School of Natural Sciences
University of California, Merced

Fall 2014

Short History of Monte Carlo

- Original Metropolis algorithm dates back to 1953 paper (J. Chem Phys) for integrals in $\mathbb{R}^{2N}$, where N denoted the number of particles in a problem of interest
- Hastings proposed a solution to the "curse of dimensionality" in 1970 paper[1] .
- Gelfand and Smith [2] (1990) brought more attention to the subject
- Advent of more computational power, especially parallel processing makes MC methods more popular
- Combination of computing power and algorithms drives new interest in Bayesian statistics

Motivating applications

Example 1: Generate samples from a given probability distribution function (pdf), $f(x)$

$$X_1, X_2, \ldots, X_t$$

Example 2: Computation of certain integrals:

$$I = \int f(x)\,\mathrm{d}x$$

Example 3: Given $f : \mathbb{R}^n \to \mathbb{R}$

$$\text{find } x_* \in \mathbb{R}^n \text{ such that } f(x_*) \leq f(x),\ \forall x \in \mathbb{R}^n$$

Some important assumptions & facts

- We will usually assume that we can compute the pdf $f(x)$ only up to a multiplicative constant
- Don't always know the normalizing constant
- Sampling from a given distribution can be quite difficult, especially in the high-dimensional case
- We should pick more samples from regions where $f(x)$ is "big", but how do we know that without knowing where f is "big"

Acceptance/Rejection Method

Given $f(x) \leq Mg(x)$
repeat
  Generate $X \sim g$ and $u \sim U(0,1)$
  if $u \leq f(X)/Mg(X)$ **then**
    $Y = X$
  end
until *(done)*;

Algorithm 1: Acceptance-Rejection (AR)

Remark

Algorithm (AR) produces a variable Y distributed according to $f(x)$.

Integration Example

Suppose we want to compute the following integral (expectation):

$$E[f(X)] = \frac{\int f(x)\pi(x)\,\mathrm{d}x}{\int \pi(x)\,\mathrm{d}x},$$

and the distribution of X is in k-dimensional Euclidian space.

First we compute samples $(X_t, t = 1, \ldots, N)$ with the given distribution $f(x)$ and then approximate

$$E[f(X)] = \frac{1}{n}\sum_{t=1}^{n} f(X_t).$$

Remark

Accuracy will depend on distribution, number of samples we generate, and dimension of the problem. The last is known as the "curse of dimensionality"

Metropolis-Hastings (1970)

Let

$$\alpha(X, Y) = \min\left(1, \frac{\pi(Y)q(X|Y)}{\pi(X)q(Y|X)}\right).$$

Initialize X_0; set $t = 0$
while *(not done)* **do**
  Generate $Y \sim q(\cdot|X_t); u \sim U(0,1)$
  if $u \leq \alpha(X_t, Y)$ **then**
    $X_{t+1} = Y$
  else
    $X_{t+1} = X_t$
  end
  Increment t
end

Algorithm 2: Metropolis-Hastings (MH)

Terminology

$q(\cdot|X)$ is called the proposal distribution.

Ising Model Part 1

Suppose we have a $2D$ array of spins $\sigma_i \in [-1, 1]$ and an energy function defined for a configuration of a system

$$E(\sigma) = -\sum_{ij} \mathcal{J}_{ij}\sigma_i\sigma_j - B\sum_k \sigma_k$$

and we want to take the mean of some function $f(\sigma)$:

$$\mathcal{F} = \frac{1}{Z(T)}\sum_\sigma f(\sigma)\exp(-E(\sigma)/\kappa T),$$

where $Z(T) = \sum_\sigma \exp(-E(\sigma)/\kappa T)$ is the *partition function*

Ising Model Part 2

Ideally would like to select configurations from the distribution

$$g(\sigma) = \frac{\exp(-E(\sigma)/\kappa T)}{Z(T)},$$

and we could approximate the mean by:

$$F = \frac{\sum_k f(\sigma_k)}{M}$$

Remark

Original Metropolis algorithm observed that if only one spin is changed, sampling g is easy. Showed that if a move lowers the energy accept it. If it raises the energy accept it with probability p. Use AR algorithm to choose site for move.

Simulated Annealing Method

Initialize X_0; set $t = 0$
while *(not converged)* **do**
  Generate new state Y and $u \sim U(0, 1)$
  if $E(Y) < E(X_t)$ **then**
    $X_{t+1} = Y$
  else
    if $u < e^{-[(E(Y)-E(X_n))/T]}$ **then**
      $X_{t+1} = Y$
    end
  end
end

Algorithm 3: Simulated Annealing (AR)

Remark

Simulated Annealing always takes a point where function is lower. But it will also accept a higher function value with a certain probability depending on the "temperature" T.

Basic Markov Chain Facts

- A Markov Chain is a sequence of random variables
- Constructed from a transition kernel K, which is a conditional probability density s.t. $X_{n+1} \sim K(X_n, X_{n+1})$
- Has stability property in that a *stationary probability distribution* exists by construction, i.e. if $X_n \sim \pi$ then $X_{n+1} \sim \pi$
- Homogeneous is when transition probability is independent of state – hence if the initial distribution or initial state is known, then the construction of MC is entirely dependent on transition

Basic Idea for Markov Chain Monte Carlo

- Use properties of Markov Chains to construct a transition kernel that has a stationary probability distribution that matches the one we want
- Assume stationarity, irreducibility, and aperiodic, we can prove Ergodic Theorem.
- Needs to also satisfy reversibility condition $\pi_i p(i \to j) = \pi_j p(j \to i)$

Practicalities

- Need to have a "burn-in" period
- How long do you run the Markov Chain
- Can use multiple chains to accelerate convergence

Summary

- Metropolis algorithm defines a Markov Chain whose limit is the desired probability distribution
- Many variations on the proposal distribution
- Growing number of applications

References I


W.K. Hastings,
Monte Carlo Sampling Methods Using Markov Chains and Their Applications, Biometrika, Vol. 57, No. 1., pp. 97-109, 1970


A.E. Gelfand and A.F.M. Smith,
Sampling-Based Approaches to Calculating Marginal Densities,
Journal of the American Statistical Association, Vol. 85, No. 410. (Jun., 1990), pp. 398-409


I. Beichi, F. Sullivan
The Metropolis Algorithm, Computing in Science and Engineering, January/February 2000.